

Lab 3: Reducing Crime (DRAFT: Stage 1)

C. Akkineni, A. Thorp, K. Hanna

November 27, 2018

Contents

Introduction (Stage 1: Draft Report)	2
Exploratory Data Analysis	2
Data Summary	2
Data Clean Up	2
Concerns about data	3
Univariate Analysis	4
Log Transforms	4
Key Interests	4
Explanatory	5
Model Analysis	11
General Crime Prediction Model	11
General Crime Prediction Model nolog crime rate	12
Simple Crime Prediction Model	13
Simple Crime Prediction Model nonlog crime rate	14
Simple Crime Prediction Model nonlog anywhere	15
Effects of Number of Police on Crime Mix	16
Preliminary Infomation (not intended to be left in)	21
From the assignment:	21
Variables:	21
Conclusion	28
Crime rate	28
Something about unfeasibility and unreliability of kitchen-sink model	28
How to reduce overall crime	28
How to reduce mix ratio (violent crimes)	28
Appendix A: Codebook	29

Introduction (Stage 1: Draft Report)

The team has been hired to provide research for a political campaign and help the campaign understand the determinants of crime and to help with policy suggestions that are applicable to local government.

```
library(knitr)
library(kableExtra)

codebook <- read.csv('codebook.csv')
crime <- read.csv('crime_v2.csv')

# Convert columns to factors and logical.
crime$county <- as.factor(crime$county)
crime$year <- as.factor(crime$year)
crime$west <- as.logical(crime$west)
crime$central <- as.logical(crime$central)
crime$urban <- as.logical(crime$urban)

# Create a log of the dependent variable
crime$logcrmrte <- log(crime$crmrte)
crime$east <- !(crime$west | crime$central)

# Possible transformation for prbconv
# crime$adjprbconv = crime$prbconv/max(crime$prbconv)

# Reorder to place logcrmrte next to crmrte and east next to central
# It's unsafe to keep this in, but handy for viewing data
# crime <- crime[,c(1,2,3,26,4:13,27,14:25)]
```

Exploratory Data Analysis

Data Summary

We were provided with a dataset of crime statistics for a selection of counties in North Carolina. After performing data clean up (outlined below) the data set contained 90 county observations each having 25 variables (outlined in the codebook found in Appendix A).

Data Clean Up

Null Rows

The dataset contained a an apostrophe 6 rows after the data which caused the csv reader to create 6 invalid rows. We feel it is safe to remove these rows as they contain no data.

```
# Delete the 6 empty observations at the end, including the row with the apostrophe.
# We can use complete.cases to do this as these 6 observations are the only incomplete observations.
crime = crime[complete.cases(crime), ]

# Fix prbconv which is a factor rather than numeric due to the apostrophe
# Convert from factor to numeric
crime$prbconv = as.numeric(as.character(crime$prbconv))
```

We found two identical observations for county 193. There is no logical reason to have two identical observations in this cross-sectional data, so we feel strongly that removing one of these two observations can only benefit our analysis.

```
# county 193 is duplicated, remove one
crime = crime[!duplicated(crime), ]
```

Concerns about data

There are three probability columns in the given dataset. Check if any of the columns has invalid values - i.e., any of the columns have less than zero or greater than 1 values.

```
#any(crime$prbarr<0 | crime$prbarr>1)
#any(crime$prbconv<0 | crime$prbconv>1)
#any(crime$prbpris<0 | crime$prbpris>1)
```

```
#summary(crime$prbarr)
#summary(crime$prbconv)
```

```
summary(crime$prbarr)[c(1,6)]
```

```
##      Min.      Max.
## 0.09277 1.09091
```

```
summary(crime$prbconv)[c(1,6)]
```

```
##      Min.      Max.
## 0.0683761 2.1212101
```

```
summary(crime$prbpris)[c(1,6)]
```

```
## Min. Max.
## 0.15 0.60
```

```
nrow(crime[(crime$prbarr<0 | crime$prbarr>1), c('county', 'prbarr')])
```

```
## [1] 1
```

```
nrow(crime[(crime$prbconv<0 | crime$prbconv>1), c('county', 'prbconv')])
```

```
## [1] 10
```

prbarr (Probability of Arrest)

We found that county 115 contained a value of 1.09 in prbarr (probability of arrest) which is not possible. It is the only observation with an invalid value, we will treat this variable as valid, however do so with caution.

prbconv (Probability of Conviction)

We found 10 observations with values greater than 1, which, again, is not a possible value for probability. We have little confidence in the veracity of this variable, and will not be performing further analysis.

Univariate Analysis

```
quick_uni_analysis = function(variable, description) {
  hist(variable, xlab = tools::toTitleCase(description),
        main = paste('Shapiro p-value:',
                     round(as.numeric(shapiro.test(variable)[2]), 6)
                     )
        )
  hist(log(crime$crmrte),
        xlab = tools::toTitleCase(paste('Log of', description)),
        main = paste('Shapiro p-value:',
                     round(as.numeric(shapiro.test(log(variable))[2]), 6)
                     )
        )
}
```

Log Transforms

(just a screen cap for now)

Var	level-level	level-log	log-log	log-level	Max
<i>prbarr</i>	0.156248664	0.223508549	1.90E-01	0.177044129	0.223509
<i>prbconv</i>	0.148969439	0.199642401	1.39E-01	0.131969092	0.199642
<i>prbpris</i>	0.002303558	0.000460971	4.85E-03	0.007255311	0.007255
<i>avgsen</i>	0.000391903	0.002437329	5.48E-04	0.004910613	0.004911
<i>polpc</i>	0.027983145	0.000108281	8.10E-02	0.151172748	0.151173
<i>density</i>	0.530523745	0.400718611	2.44E-01	0.228345246	0.530524
<i>taxpc</i>	0.201345254	0.128395649	1.15E-01	0.170334109	0.201345
<i>pctmin80</i>	0.032996936	0.054250895	1.58E-01	0.092918627	0.15784
<i>wcon</i>	0.154418782	0.155011391	1.56E-01	0.14550266	0.156377
<i>wtuc</i>	0.055693987	0.040588119	4.46E-02	0.055064172	0.055694
<i>wtrd</i>	0.182519167	0.155072453	1.52E-01	0.167400262	0.182519
<i>wfir</i>	0.112913534	0.085991251	8.31E-02	0.101587748	0.112914
<i>wser</i>	0.002711128	0.012797947	2.41E-03	0.009210632	0.012798
<i>wmfg</i>	0.124299379	0.094579197	1.25E-01	0.147665036	0.147665
<i>wfed</i>	0.240018014	0.273849011	2.51E-01	0.210404741	0.273849
<i>wsta</i>	0.039938722	0.028798795	2.29E-02	0.033344351	0.039939
<i>wloc</i>	0.129477153	0.083270789	9.17E-02	0.1330255	0.133025
<i>mix</i>	0.017424093	0.015558683	1.80E-07	0.000729444	0.017424
<i>pctymle</i>	0.084297118	0.077370015	9.72E-02	0.10508132	0.105081


Key Interests

Crimes Committed Per Person

Campaign Significance: Crime rate is a politicized and effects economy


```
quick_uni_analysis(crime$crmrte, 'crimes committed per per.')
```

Shapiro p-value: 2e-06



Crimes Committed per Per.

Shapiro p-value: 0.625962



Log of Crimes Committed per Per.

Log is preferable - both for interpretation and for better adhering to modeling assumptions


Offence Mix

This is a ratio of face-to-face crimes to all other crimes. Face-to-face crimes include violent crimes and those with a higher probability of violence, hence the more severe crimes. Focusing resources which reduces this ratio along with the overall crime rate would be more beneficial.

Campaign Significance: Violent crimes create fear and fear is a strong motivator for voters.


```
quick_uni_analysis(crime$prbarr, 'Probability of Arrest')
```

Shapiro p-value: 0



Probability of Arrest

Shapiro p-value: 0.511948




Log of Probability of Arrest

Explanatory

Probability of Arrest

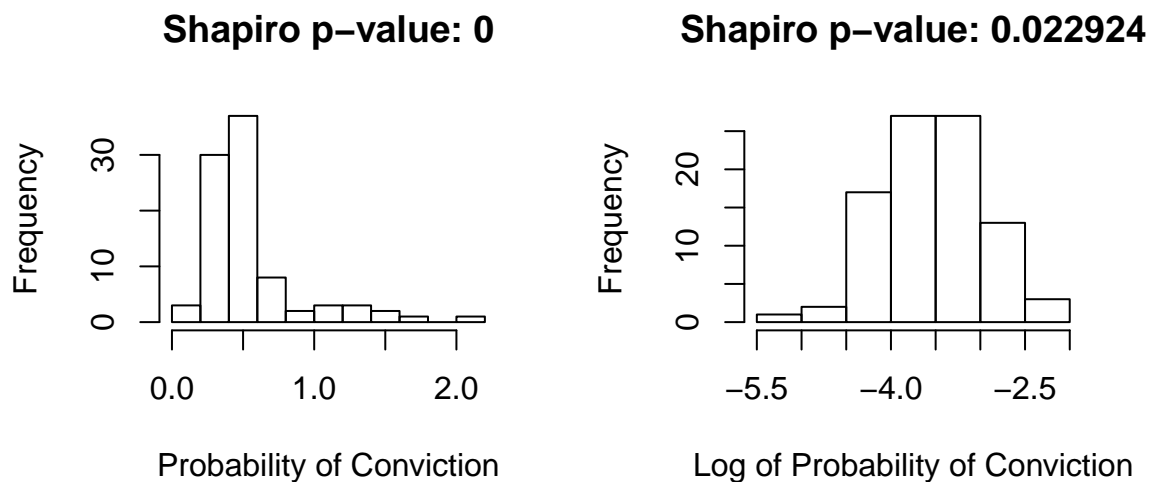
```
quick_uni_analysis(crime$prbarr, 'Probability of Arrest')
```



Log is preferable - both for interpretation and for better adhering to modeling assumptions

Probability of Conviction

```
quick_uni_analysis(crime$prbconv, 'Probability of Conviction')
```




Log is preferable - both for interpretation and for better adhering to modeling assumptions. However, even the logged version fails a Shapiro-Wilk normality test. Something to keep in mind.

Probability of Prison


```
quick_uni_analysis(crime$prbpris, 'Probability of Prison')
```

Shapiro p-value: 0.171974



Probability of Prison

Shapiro p-value: 1e-05




Log of Probability of Prison

From an interpretation standpoint, the logged version is preferable, although from an modeling assumption standpoint, the unlogged version is preferable.

Average Sentence


```
quick_uni_analysis(crime$avgsen, 'Average Sentence')
```

Shapiro p-value: 0.000464



Average Sentence

Shapiro p-value: 0.571032




Log of Average Sentence

The logged version is preferable from both an interpretation and modeling assumption standpoint.

Police as a Percentage of Population


```
quick_uni_analysis(crime$polpc, 'Police as Per. of Pop.')
```

Shapiro p-value: 0



Police as Per. of Pop.

Shapiro p-value: 1.1e-05




Log of Police as Per. of Pop.

Both logged and un-logged versions of police as a percentage of the population are non-normal. Neither is inherently preferable from a modeling assumptions standpoint.

The inclusion of this variable in a model needs to take in to consideration that an increase in crime can cause an increase in the number of police blah blah blah . Without timeseries data including changes to Police as a Percentage of Population we cannot determine the true effects. However, many studies show that increasing the size of a police force does reduce crime.


```
quick_uni_analysis(crime$pctymle, 'Per. of Pop. That Are Young Males')
```

Shapiro p-value: 0



Per. of Pop. that are Young Males

Shapiro p-value: 0




Log of Per. of Pop. that are Young Mal



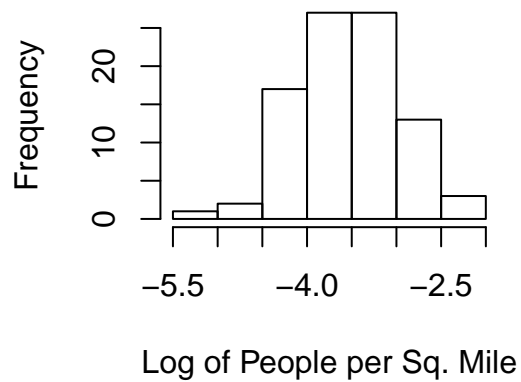
Both logged and un-logged versions of the percent of population that is young and male are non-normal. Neither is inherently preferable from a modeling assumptions standpoint.



```
quick_uni_analysis(crime$density, 'people per sq. mile')
plot(log(crime$density))
boxplot(log(crime$density))
plot(crime$logcrmrte)
boxplot(crime$logcrmrte)
```


Shapiro p-value: 0




Shapiro p-value: 0



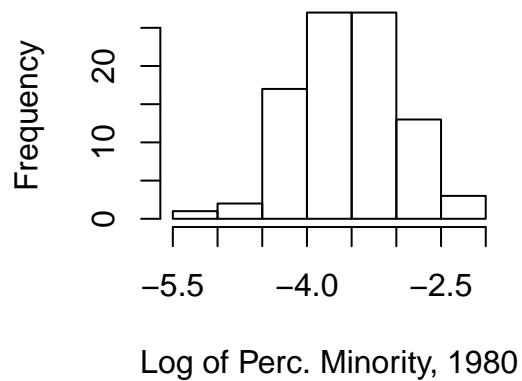


```
quick_uni_analysis(crime$pctmin80, 'perc. minority, 1980')
```

Shapiro p-value: 0.002322



Shapiro p-value: 7e-06



Model Analysis

General Crime Prediction Model

```
general_model = lm(logcrmrte ~ prbarr + prbconv + log(polpc) + log(pctmin80), data = crime)
general_aic = AIC(general_model)
general_bic = BIC(general_model)
general_rsquared = summary(general_model)[8]
general_adjrsquared = summary(general_model)[9]

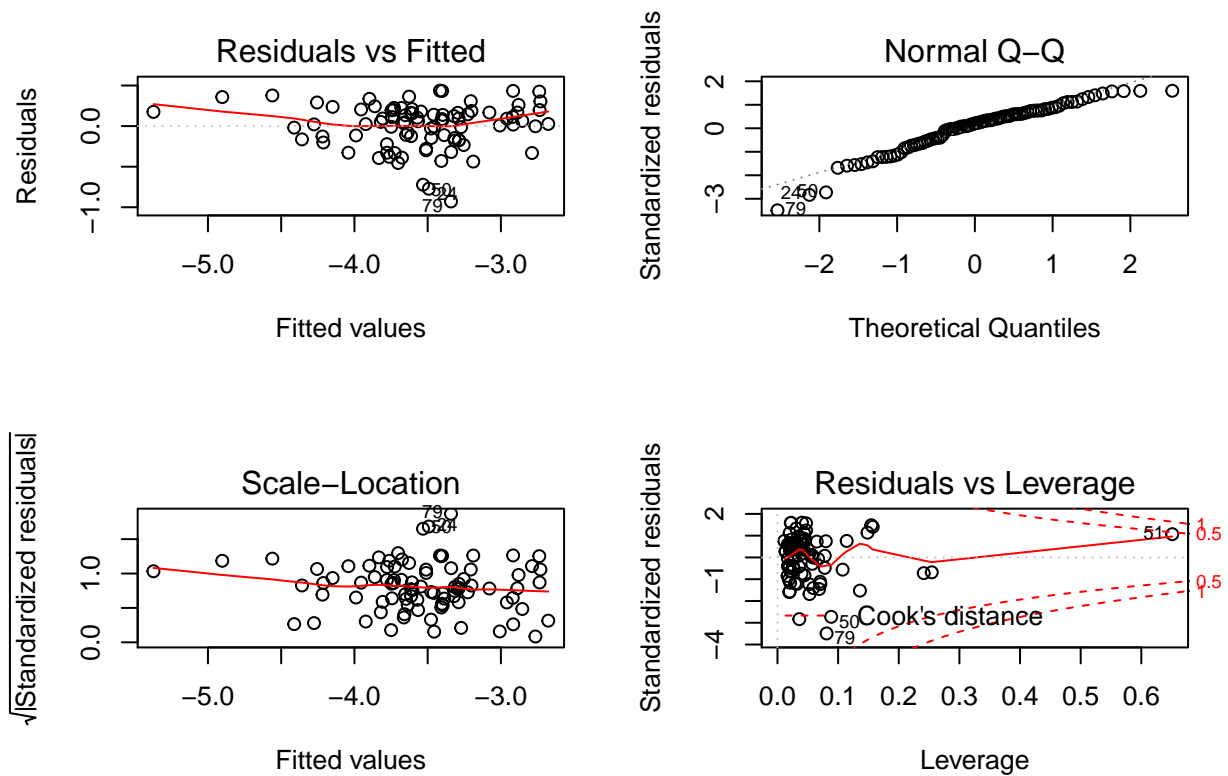
print(general_aic)

## [1] 31.44582
print(general_bic)

## [1] 46.44468
print(general_rsquared)

## $r.squared
## [1] 0.7559527
print(general_adjrsquared)

## $adj.r.squared
## [1] 0.7444682
par(mfrow = c(2,2))
plot(general_model)
```



General Crime Prediction Model nolog crime rate

```
general_model = lm(crmrte ~ log(prbarr) + log(prbconv) + polpc + log(pctmin80), data = crime)
general_aic = AIC(general_model)
general_bic = BIC(general_model)
general_rsquared = summary(general_model)[8]
general_adjrsquared = summary(general_model)[9]

print(general_aic)

## [1] -538.6185

print(general_bic)

## [1] -523.6196

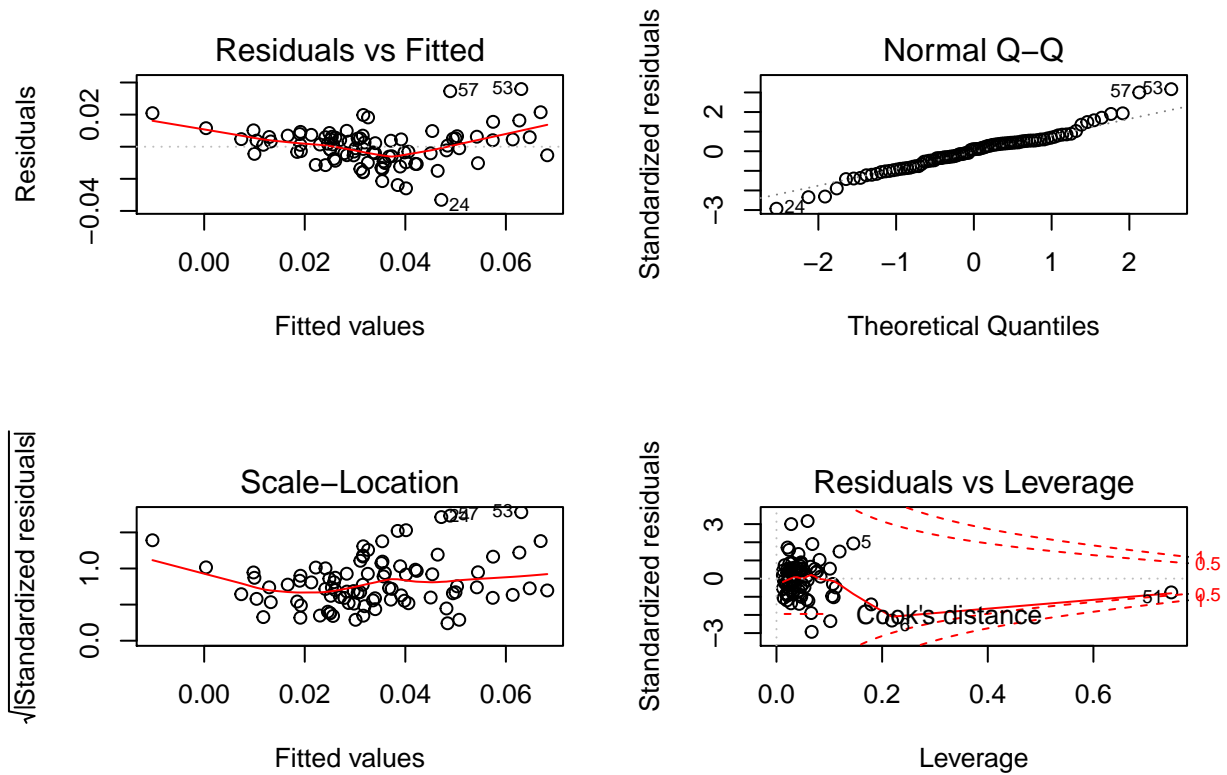
print(general_rsquared)

## $r.squared
## [1] 0.6343661

print(general_adjrsquared)

## $adj.r.squared
## [1] 0.6171598
```

```
par(mfrow = c(2,2))
plot(general_model)
```



polpc and pctmin80 are control

->

Simple Crime Prediction Model

```
general_model = lm(crmrte ~ log(prbarr) + log(prbconv), data = crime)
general_aic = AIC(general_model)
general_bic = BIC(general_model)
general_rsquared = summary(general_model)[8]
general_adjrsquared = summary(general_model)[9]

print(general_aic)
```

```
## [1] -496.07
```

```
print(general_bic)
```

```
## [1] -486.0707
```


```
print(general_rsquared)
```

```
## $r.squared
## [1] 0.3867098
```

```
print(general_adjrsquared)
```

```
## $adj.r.squared
## [1] 0.3726112
```

```
par(mfrow = c(2,2))
plot(general_model)
```



Simple Crime Prediction Model nonlog crime rate

```
general_model = lm(crmrte ~ log(prbarr) + log(prbconv) + polpc + log(pctmin80), data = crime)
general_aic = AIC(general_model)
general_bic = BIC(general_model)
general_rsquared = summary(general_model)[8]
general_adjrsquared = summary(general_model)[9]
```

```
print(general_aic)
```

```
## [1] -538.6185
```

```
print(general_bic)
```

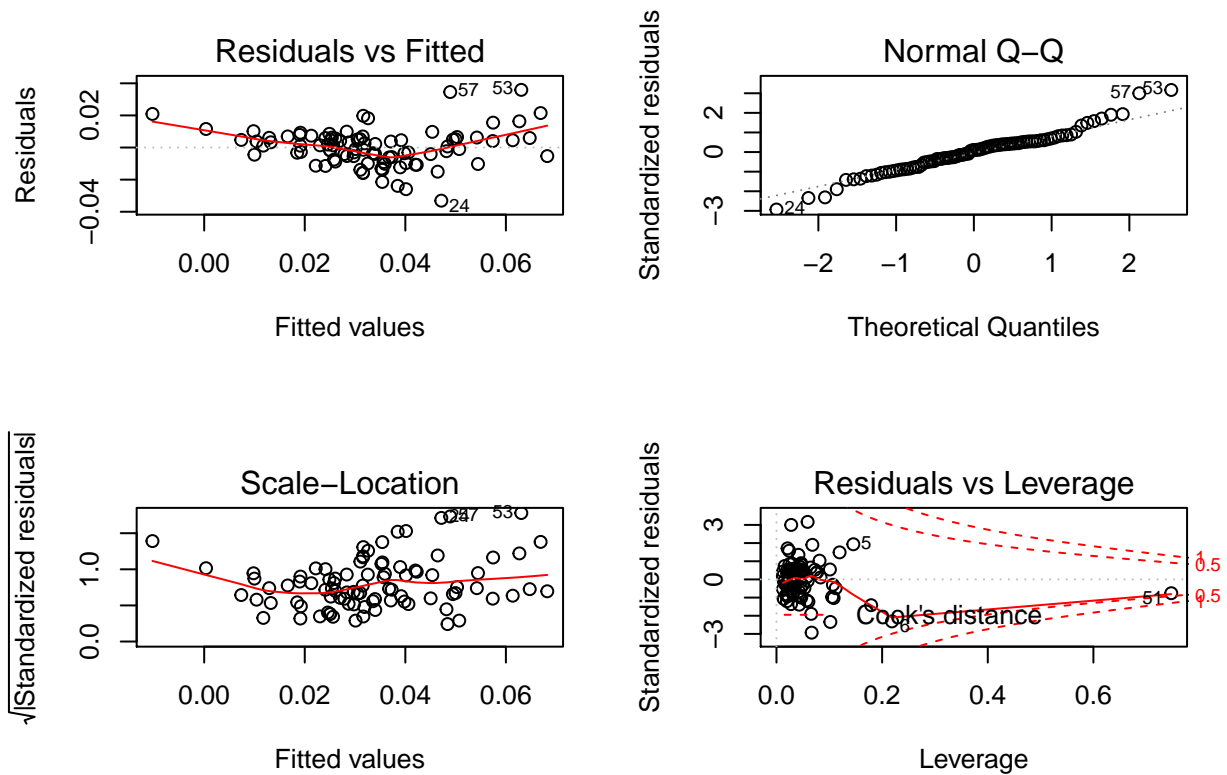
```
## [1] -523.6196
```

```
print(general_rsquared)
```

```
## $r.squared
```

```
## [1] 0.6343661
print(general_adjrsquared)

## $adj.r.squared
## [1] 0.6171598
par(mfrow = c(2,2))
plot(general_model)
```



Simple Crime Prediction Model nonlog anywhere

```
general_model = lm(crmrte ~ prbarr + prbconv + polpc + pctmin80, data = crime)
general_aic = AIC(general_model)
general_bic = BIC(general_model)
general_rsquared = summary(general_model)[8]
general_adjrsquared = summary(general_model)[9]

print(general_aic)

## [1] -546.4089
print(general_bic)

## [1] -531.4101
```


```
print(general_rsquared)
```

```
## $r.squared
## [1] 0.6646843
```

```
print(general_adjrsquared)
```

```
## $adj.r.squared
## [1] 0.6489048
```

```
par(mfrow = c(2,2))
plot(general_model)
```



->

Effects of Number of Police on Crime Mix

```
general_model = lm(mix ~ polpc + pctmin80 + log(100*density) + pctymle, data = crime)
general_aic = AIC(general_model)
general_bic = BIC(general_model)
general_rsquared = summary(general_model)[8]
general_adjrsquared = summary(general_model)[9]

print(general_aic)
```

```
## [1] -207.0104
```



```
print(general_bic)
```

```
## [1] -192.0116
```

```
print(general_rsquared)
```

```
## $r.squared
```

```
## [1] 0.2229964
```

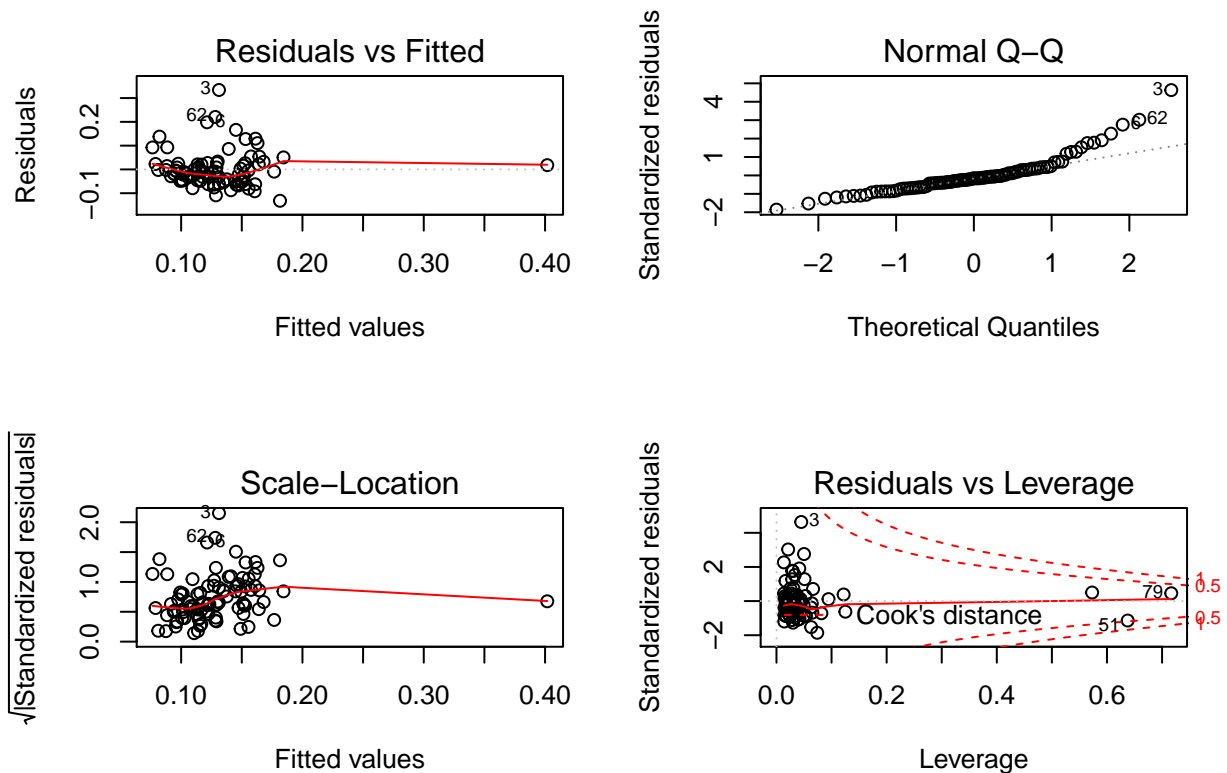
```
print(general_adjrsquared)
```

```
## $adj.r.squared
```

```
## [1] 0.1864315
```

```
par(mfrow = c(2,2))
```

```
plot(general_model)
```



```
summary(general_model)
```

```
##
```

```
## Call:
```

```
## lm(formula = mix ~ polpc + pctmin80 + log(100 * density) + pctymle,
```

```
## data = crime)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
```

```
## -0.13180 -0.04516 -0.01160  0.02103  0.33391
```

```
##
```

```
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.2262385  0.0424190   5.333 7.84e-07 ***
## polpc          2.4570101  8.0405300   0.306  0.7607
## pctmin80       0.0007944  0.0004697   1.691  0.0945 .
## log(100 * density) -0.0250531  0.0057998  -4.320 4.22e-05 ***
## pctymle        -0.1022449  0.3376326  -0.303  0.7628
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.07375 on 85 degrees of freedom
## Multiple R-squared:  0.223, Adjusted R-squared:  0.1864
## F-statistic: 6.099 on 4 and 85 DF,  p-value: 0.0002308
```

Police in High Density

```
general_model = lm(crmrte ~ polpc, data = crime)
general_aic = AIC(general_model)
general_bic = BIC(general_model)
general_rsquared = summary(general_model)[8]
general_adjrsquared = summary(general_model)[9]

print(general_aic)

## [1] -456.6218

print(general_bic)

## [1] -449.1224


print(general_rsquared)

## $r.squared
## [1] 0.02798314

print(general_adjrsquared)

## $adj.r.squared
## [1] 0.0169375

par(mfrow = c(2,2))
plot(general_model)
```



```
summary(general_model)
```

```
##
## Call:
## lm(formula = crmrte ~ polpc, data = crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.051400 -0.011799 -0.003837  0.006455  0.063787
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.02806     0.00395   7.105 2.99e-10 ***
## polpc        3.18839     2.00318   1.592  0.115
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.01873 on 88 degrees of freedom
## Multiple R-squared:  0.02798,    Adjusted R-squared:  0.01694
## F-statistic: 2.533 on 1 and 88 DF,  p-value: 0.115

general_model = lm(crmrte ~ polpc, data = crime[crime$density >= quantile(crime$density, .90), ])
general_aic = AIC(general_model)
general_bic = BIC(general_model)
general_rsquared = summary(general_model)[8]
general_adjrsquared = summary(general_model)[9]
```

```
print(general_aic)
```

```
## [1] -42.11032
```

```
print(general_bic)
```

```
## [1] -41.51865
```

```
print(general_rsquared)
```

```
## $r.squared
```

```
## [1] 0.09167394
```


```
print(general_adjrsquared)
```

```
## $adj.r.squared
```

```
## [1] -0.03808692
```

```
par(mfrow = c(2,2))
```

```
plot(general_model)
```



```
summary(general_model)
```

```
##
```

```
## Call:
```

```
## lm(formula = crmrte ~ polpc, data = crime[crime$density >= quantile(crime$density,  
##   0.9), ])
```

```
##
```

```
## Residuals:
```

```
##           Min           1Q       Median           3Q           Max
## -0.020053 -0.010839 -0.006470  0.006215  0.029749
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.01420    0.06478   0.219   0.833
## polpc       24.65711    29.33533   0.841   0.428
##
## Residual standard error: 0.01895 on 7 degrees of freedom
## Multiple R-squared:  0.09167,    Adjusted R-squared:  -0.03809
## F-statistic: 0.7065 on 1 and 7 DF,  p-value: 0.4284
```

More police in higher density correlate to higher crime rate. Probably not worth keeping, but interesting.

Preliminary Information (not intended to be left in)

From the assignment:

- 1. What do you want to measure? Make sure you identify variables that will be relevant to the concerns of the political campaign.
- 2. What transformations should you apply to each variable? This is very important because transformations can reveal linearities in the data, make our results relevant, or help us meet model assumptions.
- 3. Are your choices supported by EDA? You will likely start with some general EDA to detect anomalies (missing values, top-coded variables, etc.). From then on, your EDA should be interspersed with your model building. Use visual tools to guide your decisions.
- 4. What covariates help you identify a causal effect? What covariates are problematic, either due to multicollinearity, or because they will absorb some of a causal effect you want to measure?

Variables:

1. Target
 - crmrte
2. Label
 - county
3. Geographic:
 - density (likely related to others, especially urban)
 - west
 - central
 - urban
4. Certainty of Punishment:
 - prbarr
 - prbconv

- polpc

5. Severity of Punishment:

- prbpris
- avgscn

We haven't really addressed what kinds of crimes are being talked about. Should we assume that it is violent crimes and drug crimes? (i.e. not white collar or other non-violent) If we also assume that some crimes are committed by rational actors weighing gains versus losses, we would expect crimes like theft and drug dealing to be negatively correlated with wealth (i.e. higher opportunity cost for wealthy people to engage in those sorts of crimes - and by extension linked crimes like manslaughter (I'm thinking an attempted theft or drug deal gone wrong).

The idea here is that higher real wages reduce crime propensity. To operationalize that, ideally we would have some idea of cost of living and nominal wages (particularly in sectors that are potential alternatives to those who typically engage in violent crime and drug dealing - those without high educational barriers to entry - construction, driving, retail).

In addition to pure wealth, there are other ideas like social capital and a sense of belonging/rootedness that we can imagine would be associated with lower propensity to commit crimes. We could assess this through variables like homeownership, rate of volunteering, rates of religious attendance, or the proportion of mothers that are single and have sole custody.

Another thought is that educated people are also less likely to commit these kinds of crimes. Another relevant variable may be average years of education, proportion of population with at least a high school diploma, at least a bachelor's, or one or more graduate degrees. Related to this category as well as the previous notion of social belonging, the rate of absenteeism among high school students could also reflect these two notions.

Another idea is that proximity to instruments of crime also increases propensity to commit crimes. The rate of gun ownership, or - even better - the rate of illicit gun ownership would be useful here.

Another idea is that crimes are more likely as inequality is high between neighboring locations - creating a threshold upon which economic and extension crimes become focused. This is hard to assess without better location data

And another idea is that organized crime tends to be more persistent than isolated instances of crime. So, to the extent possible, it would be useful to have a measure of gang/mob/etc membership as a percent of the population.

Recidivism would be interesting, how much is caused by repeat offenders, how sentencing affects overall crime.

```
mod1 <- lm(crmrte ~ prbarr + prbconv + prbpris + avgscn, data=crime)
summary(mod1)
```

```
##
## Call:
## lm(formula = crmrte ~ prbarr + prbconv + prbpris + avgscn, data = crime)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-0.031434	-0.009597	-0.002289	0.007937	0.052989

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.0446704	0.0109773	4.069	0.000105 ***
prbarr	-0.0625420	0.0121899	-5.131	1.80e-06 ***
prbconv	-0.0234739	0.0047115	-4.982	3.27e-06 ***
prbpris	0.0212198	0.0204908	1.036	0.303336

```
## avgsen      0.0011892  0.0006007   1.980 0.050992 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.01549 on 85 degrees of freedom
## Multiple R-squared:  0.358, Adjusted R-squared:  0.3278
## F-statistic: 11.85 on 4 and 85 DF,  p-value: 1.072e-07
mod1_log <- lm(log(crmrte) ~ log(prbarr) + log(prbconv) + log(prbpris) + log(avgsen), data=crime)
summary(mod1_log)
```

```
##
## Call:
## lm(formula = log(crmrte) ~ log(prbarr) + log(prbconv) + log(prbpris) +
##     log(avgsen), data = crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.36104 -0.19129  0.07939  0.27754  0.86843
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -4.86792     0.43153  -11.281  < 2e-16 ***
## log(prbarr)   -0.72397     0.11532   -6.278 1.39e-08 ***
## log(prbconv) -0.47251     0.08311   -5.686 1.80e-07 ***
## log(prbpris)  0.15967     0.20644    0.773  0.441
## log(avgsen)   0.07642     0.16347    0.467  0.641
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.429 on 85 degrees of freedom
## Multiple R-squared:  0.4162, Adjusted R-squared:  0.3888
## F-statistic: 15.15 on 4 and 85 DF,  p-value: 2.171e-09
```

```
mod2 <- lm(crmrte ~ prbarr + prbconv + prbpris + avgsen + polpc, data=crime)
summary(mod2)
```

```
##
## Call:
## lm(formula = crmrte ~ prbarr + prbconv + prbpris + avgsen + polpc,
##     data = crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.038338 -0.007831 -0.000843  0.006578  0.039503
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.0551317  0.0094088   5.860 8.80e-08 ***
## prbarr       -0.0896980  0.0112238  -7.992 6.30e-12 ***
## prbconv      -0.0272321  0.0040173  -6.779 1.57e-09 ***
## prbpris       0.0121931  0.0173230   0.704  0.483
## avgsen       -0.0003331  0.0005662  -0.588  0.558
## polpc        10.5869888  1.7684871   5.986 5.11e-08 ***
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.01304 on 84 degrees of freedom
## Multiple R-squared:  0.55, Adjusted R-squared:  0.5232
## F-statistic: 20.53 on 5 and 84 DF,  p-value: 2.447e-13
mod2_log <- lm(log(crmrte) ~ log(prbarr) + log(prbconv) + log(prbpris) + log(avgsen) + log(polpc), data=crime)
summary(mod2_log)

##
## Call:
## lm(formula = log(crmrte) ~ log(prbarr) + log(prbconv) + log(prbpris) +
##     log(avgsen) + log(polpc), data = crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.40193 -0.22316  0.07635  0.25674  0.69986
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -1.70812     1.05253  -1.623  0.10836
## log(prbarr)   -0.73391     0.10934  -6.712 2.12e-09 ***
## log(prbconv) -0.43409     0.07965  -5.450 4.94e-07 ***
## log(prbpris)  0.13076     0.19587   0.668  0.50623
## log(avgsen)  -0.14564     0.16927  -0.860  0.39201
## log(polpc)    0.41412     0.12710   3.258  0.00162 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4066 on 84 degrees of freedom
## Multiple R-squared:  0.4817, Adjusted R-squared:  0.4509
## F-statistic: 15.62 on 5 and 84 DF,  p-value: 7.621e-11
```

I disagree with the inclusion of police percentage in this regression. While all else equal, we may expect more police to result in a deterrent, we can only assess this against an unobservable counterfactual - how the same location would have been impacted had more police or less police been assigned there. This would be possible to address experimentally, but not observationally. The problem is that places with increasing crime rates are subsequently likely to increase police presence, while those with decreasing crime rates are likely to decrease police presence (for budgetary reasons).

- I found studies showing increasing police decreases crime, including studies done with accidental experiments. I cited a couple above in the univariate (only visible in code because of HTML comments)

```
mod3 <- lm(crmrte ~ prbarr + prbconv + prbpris + avgsen + density, data=crime)
summary(mod3)
```

```
##
## Call:
## lm(formula = crmrte ~ prbarr + prbconv + prbpris + avgsen + density,
##     data = crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.019774 -0.008035 -0.002675  0.005583  0.043581
##
## Coefficients:
```



```
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.0347037  0.0084464   4.109 9.21e-05 ***
## prbarr       -0.0338981  0.0099540  -3.405 0.001015 **
## prbconv      -0.0148290  0.0037470  -3.958 0.000158 ***
## prbpris       0.0055122  0.0157162   0.351 0.726668
## avgssen       0.0004454  0.0004666   0.954 0.342582
## density       0.0072471  0.0009143   7.926 8.52e-12 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.01178 on 84 degrees of freedom
## Multiple R-squared:  0.6327, Adjusted R-squared:  0.6108
## F-statistic: 28.94 on 5 and 84 DF,  p-value: < 2.2e-16
mod3_log <- lm(log(crmrte) ~ log(prbarr) + log(prbconv) + log(prbpris) + log(avgssen) + log(density), data=crime)
summary(mod3_log)
```

```
##
## Call:
## lm(formula = log(crmrte) ~ log(prbarr) + log(prbconv) + log(prbpris) +
##      log(avgssen) + log(density), data = crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.17789 -0.21695  0.06725  0.25095  0.80522
##
## Coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -4.57783     0.40718 -11.243 < 2e-16 ***
## log(prbarr)   -0.53503     0.11764  -4.548 1.81e-05 ***
## log(prbconv)  -0.41601     0.07845  -5.303 9.07e-07 ***
## log(prbpris)  -0.23672     0.21732  -1.089 0.279155
## log(avgssen)  -0.08037     0.15695  -0.512 0.609946
## log(density)   0.14780     0.03835   3.854 0.000227 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.3978 on 84 degrees of freedom
## Multiple R-squared:  0.5039, Adjusted R-squared:  0.4744
## F-statistic: 17.07 on 5 and 84 DF,  p-value: 1.292e-11
```

I also disagree with this simplistic inclusion of density (if our goal is causation, not prediction). All else equal - density means more eyes looking around - which should mean less propensity to commit crime. There are likely other variables which better explain what many lazily associate with density. For example, thresholds of inequality, where extreme wealth exists right next to extreme poverty are more common in more dense locations. Additionally, more dense locations sometimes have less average social capital (perhaps again because of close proximity class differences) than less dense places. These are good instances of omitted variable bias. Since our goal is causation rather than prediction, we should be cautious in including density as an explanatory variable.

- Interesting, the last line in in the abstract here: <https://www.ncjrs.gov/App/Publications/abstract.aspx?ID=99314> agrees. Correlation, not likely causation.

```
mod4 <- lm(crmrte ~ prbarr + prbconv + prbpris + avgssen + pctymle, data=crime)
summary(mod4)
```

```
##
## Call:
## lm(formula = crmrte ~ prbarr + prbconv + prbpris + avgse + pctymle,
##     data = crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.031869 -0.008483 -0.002073  0.006987  0.053989
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.0333499   0.0129958   2.566  0.0121 *
## prbarr       -0.0583518   0.0123644  -4.719 9.35e-06 ***
## prbconv      -0.0219883   0.0047617  -4.618 1.38e-05 ***
## prbpris       0.0231725   0.0203451   1.139  0.2580
## avgse        0.0010613   0.0006008   1.767  0.0809 .
## pctymle      0.1154660   0.0725020   1.593  0.1150
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.01535 on 84 degrees of freedom
## Multiple R-squared:  0.3768, Adjusted R-squared:  0.3397
## F-statistic: 10.16 on 5 and 84 DF,  p-value: 1.24e-07
mod4_log <- lm(log(crmrte) ~ log(prbarr) + log(prbconv) + log(prbpris) + log(avgse) + log(pctymle), da
summary(mod4_log)

##
## Call:
## lm(formula = log(crmrte) ~ log(prbarr) + log(prbconv) + log(prbpris) +
##     log(avgse) + log(pctymle), data = crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.3599 -0.1852  0.0841  0.2844  0.8779
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -4.35261    0.86966  -5.005 3.03e-06 ***
## log(prbarr)   -0.69480    0.12331  -5.635 2.29e-07 ***
## log(prbconv) -0.45446    0.08746  -5.196 1.40e-06 ***
## log(prbpris)  0.16249    0.20714   0.784  0.435
## log(avgse)    0.06348    0.16508   0.385  0.702
## log(pctymle)  0.17282    0.25296   0.683  0.496
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4304 on 84 degrees of freedom
## Multiple R-squared:  0.4195, Adjusted R-squared:  0.3849
## F-statistic: 12.14 on 5 and 84 DF,  p-value: 7.347e-09
```

The inclusion of the percent of the population that is young and male, as expected, has a notable positive relationship with crime rate. The variable is only marginally significant though. I would argue that, since we are interested in causation, particularly in differentiating the effect of certainty of punishment from severity of punishment, we should include this variable so as to make the coefficients for certainty and severity of

punishment less biased.

The logged version of this regression is not anywhere close to significant for the log(pctymle) variable.

Interesting things to note after examining these regression approaches:

- 1) The certainty of punishment is estimated much more consistently than any other variable. In every regression run, these variables have a clear, significant negative relationship with crime rate. This is a good basis upon which we can make policy recommendations - Our policy recommendations focus on increasing the certainty of punishment in areas with low certainty of punishment now. Police departments need quality information in order to increase the certainty of punishment, so improving relationships with local people is a key recommendation. Additionally, locations with a low conviction rate could invest in more prosecutorial staff/better prosecutorial staff.
- 2) The severity of punishment seems to be much less consistent, with the sign flipping depending on the regression specification. In contrast to some people's expectations, most regressions show that the probability of going to prison is associated with a higher crime rate, rather than a lower one. One way of thinking about this is that prisons have become cultivators of more organized crime. Another way to think about it is prisons currently do little to change the person put in them for the better. Rather, prisoners associate all day every day with other prisoners - perhaps exchanging bad habits and temperaments and building social connections that may persist in and outside of prison.
- 3) Police presence and density are both useful predictors, though we believe that they actually obscure the causal relationship we desire to study. Police presence tends to increase along with crime and decrease as crime goes away. In other words, the direction of causality is likely reversed from what we have specified here. Secondly, density is a useful predictor, but we feel it does not contain causal information about crime rates. Rather density obscures more relevant but unobserved relationships with other factors such as proximal inequality.

```
crime$urban + crime$west + crime$central
```

```
## [1] 1 1 1 1 1 1 0 0 0 0 2 1 1 1 1 1 1 0 1 0 0 1 0 0 1 1 0 2 0 2 1 2 1 0
## [36] 2 0 0 1 1 0 0 1 1 0 1 0 1 1 1 1 0 2 1 1 0 1 0 0 1 0 0 0 0 1 0 1 1 1 0
## [71] 1 1 1 0 0 1 1 1 1 1 1 1 2 1 0 1 0 1 0 1
```

```
crime$urban + crime$west
```

```
## [1] 0 0 1 0 1 1 0 0 0 0 2 1 0 1 0 0 0 1 0 0 0 0 1 0 0 0 0 0 1 0 1 0 1 0 0
## [36] 1 0 0 1 1 0 0 0 1 0 0 0 0 1 1 1 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 1 0 0 0
## [71] 0 0 1 0 0 0 0 1 1 1 0 0 1 0 0 1 0 1 0 1
```

```
crime$urban + crime$central
```

```
## [1] 1 1 0 1 0 0 0 0 0 0 1 0 1 0 1 1 1 0 0 1 0 0 1 0 0 1 1 0 2 0 2 1 1 1 0
## [36] 2 0 0 0 0 0 0 1 0 0 1 0 1 0 0 0 0 2 1 1 0 1 0 0 1 0 0 0 0 1 0 0 1 1 0
## [71] 1 1 0 0 0 1 1 0 0 0 1 1 2 1 0 0 0 0 0 0
```

```
crime$west + crime$central
```

```
## [1] 1 1 1 1 1 1 0 0 0 0 1 1 1 1 1 1 1 1 0 1 0 0 0 0 0 1 1 0 1 0 1 1 2 1 0
## [36] 1 0 0 1 1 0 0 1 1 0 1 0 1 1 1 1 0 1 1 1 0 0 0 0 1 0 0 0 0 1 0 1 1 1 0
## [71] 1 1 1 0 0 1 1 1 1 1 1 1 1 1 0 1 0 1 0 1
```

Conclusion

From the assignment “Since you are restricted to ordinary least squares regression, omitted variables will be a major obstacle to your estimates. You should aim for causal estimates, while clearly explaining how you think omitted variables may affect your conclusions.”

Crime rate

Can use crime rate as leverage against incumbent if it's high, or to defend as the incumbent if it's low.

Something about unfeasibility and unreliability of kitchen-sink model

While adding police, increasing sentences, pay everybody more money,

How to reduce overall crime

From selected models

How to reduce mix ratio (violent crimes)

More police reduces mix ratio

Appendix A: Codebook

Table 1: Crime Data Codebook

Variable	Label	Notes
county	county identifier	
year	1987	
crmrte	crimes committed per person	
prbarr	'probability' of arrest	County 115 has a value of 1.09, which is not a possible probability. There are 10 observations greater than 1, which is not a possible probability.
prbconv	'probability' of conviction	
prbpris	'probability' of prison sentence	
avgsen	avg. sentence, days	
polpc	police per capita	
density	people per sq. mile	
taxpc	tax revenue per capita	
west	=1 if in western N.C.	
central	=1 if in central N.C.	
urban	=1 if in SMSA	
pctmin80	perc. Minority, 1980	
wcon	weekly wage, construction	
wtuc	wkly wge, trns, util, commun	
wtrd	wkly wge whlesle, retail, trade	
wfir	wkly wge, fin, ins, real est	
wser	wkly wge, service industry	
wmfg	wkly wge, manufacturing	
wfed	wkly wge fed employees	
wsta	wkly wge state employees	
wloc	wkly wge local gov emps	
mix	offense mix: face-to-face/other	
pctymle	percent young male	